More relevant music.
Not a smarter algorithm.
Just better memory.

local-mood is an MCP server that builds Spotify playlists. The idea is simple: better memory leads to better outcomes. This demo proves it with your own listening history.

Model quality gets the headlines. Memory decides the outcomes.

Why this tool exists:

✦ Joined Cronos Labs, Dec 2025

Conversations lived in DMs; decisions weren't written down. The missing context showed up in product velocity and quality.

THE PROBLEM

✦ Rebuilt the info architecture

With IT: standard project channels, shared team drives, mandated artifacts across the product process. That corpus feeds our company brain.

THE FIX

✦ Built this demo

Even the best process requires change management. I wanted the team to feel it, not take my word for it, so the proof is their own music.

WHY A SPOTIFY TOOL

Better memory → better outcomes.

Spotify's API remembers 50 plays. Spotify remembers everything.

A new app gets ~50 recent plays plus three rank-only top-track lists. No play counts, no timestamps, no history. Spotify also removed the recommendations endpoint for new apps, so similar-tracks is off the table. But every user can export their full streaming history, and mine goes back to 2019.

50

PLAYS THE API SHOWS

A context window. Recency and rank, nothing else.

56,784

STREAMS MY EXPORT REMEMBERS

Every stream since 2019. 30,799 unique tracks, 7.3 years.

1,135×

THE MEMORY MULTIPLIER

Same account, same music, same API scopes. Only memory differs.

Every AI system has this exact gap.

This is the same gap every AI product has. Music just makes it obvious.

✦ The API window

What the system can see right now. Recency and rank, no explanation. Ask it why a song is a favorite and the best it can say is "it's in your top tracks."

= A CONTEXT WINDOW

✦ The export

Play counts, completions, skips, deliberate starts, hour-of-day, loyalty across years. Actual behavior, with signals no cleverness can reconstruct from a 50-play window.

= LONG-TERM MEMORY

✦ The labels

Claude listened through my library once and labeled 403 tracks with emotions. Written down, "melancholy" becomes a query instead of a guess made fresh every session.

= SEMANTIC MEMORY

Better memory → better outcomes.

A memory and scoring layer between Claude and Spotify.

Claude + Spotify's official connector

Claude
asks for a playlist

↓

official Spotify connector
sends the ask to Spotify

↓

Spotify's recommender
runs on their servers, on the listening data they keep

↓

playlist
good picks, no way to see why

Spotify picks the songs on their own servers, using the listening data they keep. You get good picks, but you can't see that data, ask why a song was chosen, or build on it.

I tested it: the same throwback ask, run twice, returned two different playlists. Checked against my export, 13 of its 25 picks are songs I've played fewer than 10 times ever.

Developers building their own apps don't get this engine at all. They get the last 50 plays, and that's the path the next two slides measure.

Claude + local-mood

Claude
asks for a playlist

↓

LOCAL-MOOD MCP · ONE INSTALL, THE SPOTIFY API WRAPPED INSIDE

memory database
56,784 streams · every sync journaled · 403 labeled tracks

↓

scoring engine
completion · skips · loyalty · era: deterministic picks, not guesses

↓

Spotify API client
search · create · play

↓

playlist
picked from 7.3 years, with a why per track

Claude interprets the ask. The layer remembers every stream, computes the signals, and makes the picks. The wrapped API just executes.

Both apps can answer this. One answers better.

`compare_memory("throwback", count=15)`

Throwback: older tracks you still hold onto. One algorithm, run with and without memory. Both are real playlists on my Spotify.

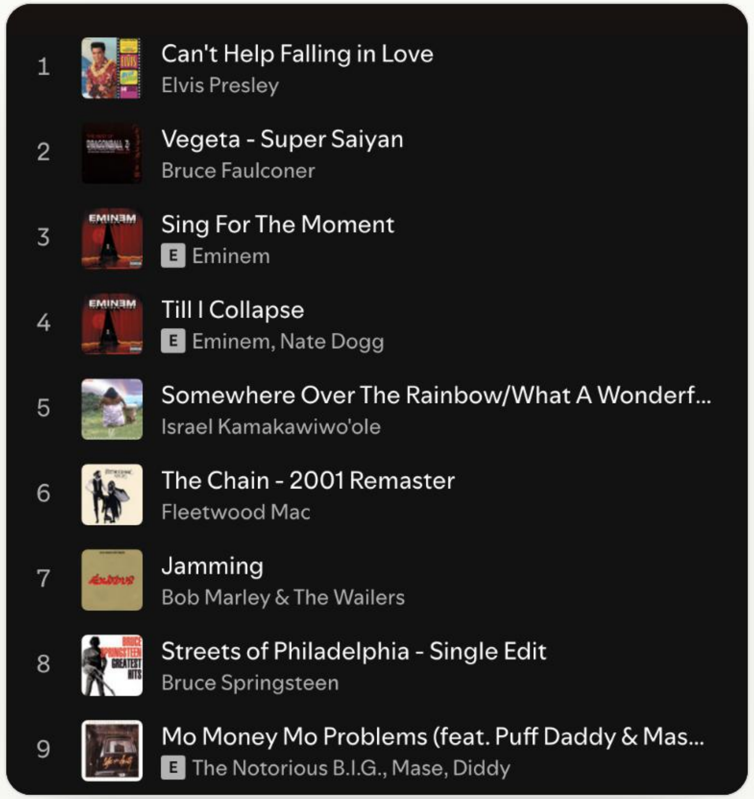

Run 1 · API only

It saw me play Sweet Caroline and Elvis at one party and decided that's who I am. Two to five plays each, no history behind any of it.

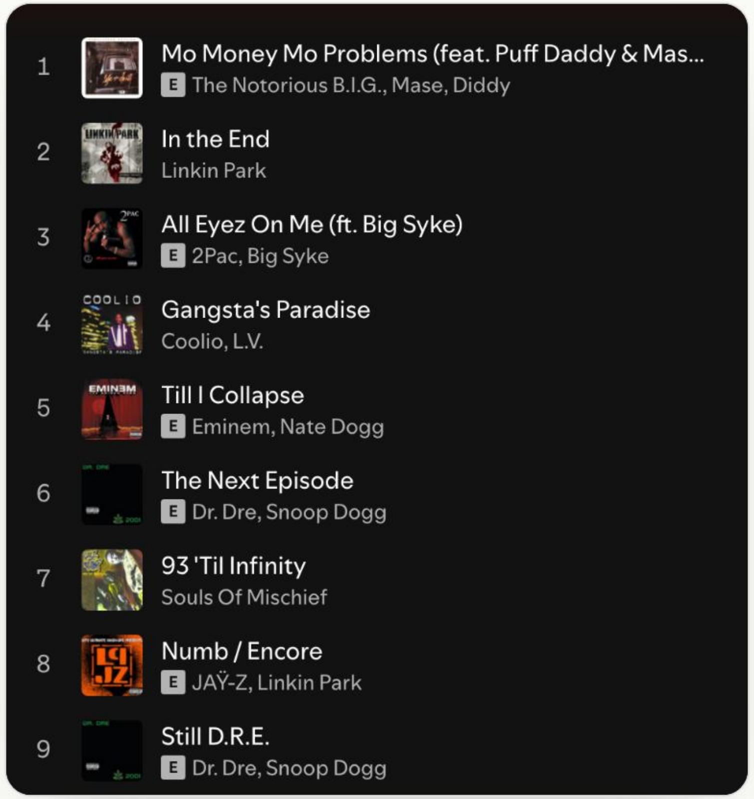

Run 2 · with memory

What I actually grew up on: Biggie, Pac, Dre, Linkin Park. 13 of 15 picks changed. Numb / Encore has 322 lifetime plays; the API can't see one of them.

This ask, the API **can't answer at all.**

`compare_memory("comfort", count=15)`

Comfort: long-loved tracks you finish and rarely skip.

```
// everything the official Web API
// gives a new app
top_tracks        rank order only
recently_played   last 50 streams
saved_tracks      yes / no

// what "comfort" is made of
completion ratio  not in the API
skip rate         not in the API
years of loyalty  not in the API
```

Run 1 · API only

The inputs don't exist, so local-mood refuses to fake it: 0 of 15 picks reproducible from the window.

1	Yours	SG Lewis
2	Before Paris	Tom Misch
3	It Runs Through Me	Tom Misch, De La Soul
4	Sad People	E Kid Cudi
5	Know Myself	E Justine Skye, Vory
6	On My Mind (Jorja Smith X Preditah)	Jorja Smith, Preditah
7	White	Frank Ocean, John Mayer
8	Thinkin Bout You	Frank Ocean
9	Swim Good	Frank Ocean

Run 2 · with memory

A real playlist on my Spotify. Every pick is backed by completion, skip, and loyalty numbers from 7.3 years of streams.

Music that fits the day I've had. One sentence to get it.

The outcome I wanted: the mood I'm in, the ramp I need next, and a few songs I've never heard. Memory does the work.

ME Given the day I've had across Telegram and Slack, figure out the mood I'm in and build a playlist that walks me from there up to leg-day energy. Sneak in a few songs I've never heard.

CLAUDE
```
telegram · get_briefing("today") ✓ · slack · recap ✓ read: long day, two escalations

mood read: drained · target: energetic, motivated

generate_playlist ✓ 20 tracks, ordered calm to energetic from my semantic labels

discover ✓ my database says energetic for me is 2000s hip-hop, so it runs a Spotify catalog search: genre "hip hop", 2000–2009

✓ ~100 candidates back from Spotify, checked one by one against my database: 3 have never been played, those get mixed in

play(device="Zain's iPhone") ✓ playback started
```

No recommendation engine anywhere. One Spotify search, plus a database that knows every song I've played.

Cronos Labs now runs on this loop.

Capture, persist, recall. This is how we work today.

✦ Decisions, written down

P/PS docs and PRDs capture the why behind every call, once. Teams stop re-litigating last quarter's decisions and spend that energy shipping. Decision memory is what makes "less deliberating" possible.

DECISION MEMORY

✦ Retros write learnings

Our content machine runs weekly engagement retros that write what worked back into per-channel learnings files. Every campaign draft starts from everything the last one taught us instead of from zero.

LEARNING LOOPS

✦ Context, pre-computed

A daily job drops the overnight Telegram briefing into the workspace before anyone sits down. People and agents start the day already caught up, instead of spending the first meeting rebuilding context.

STANDING CONTEXT

Better memory → better outcomes.

Clone it. Drop in your export. Hear the difference.

MIT licensed. Works with any MCP client. Request your Spotify export today, drop the files in, and run the same comparison on your own history.

`github.com/zainbacchus/local-mood-mcp`