

More relevant music. Not a smarter algorithm. Just better memory.

local-mood is an MCP server that builds Spotify playlists. The idea is simple: better memory leads to better outcomes. This demo proves it with your own listening history.

Model quality gets the headlines. Memory decides the outcomes.



I'm Zain, when I joined Cronos Labs as Chief Product Officer in December 2025, I saw our conversations lived in DMs and decisions weren't written down anywhere. The cost of this missing context was visible in product velocity and quality.

So during those first weeks, I worked in tandem with our IT team to overhaul our information architecture and data retention policies to improve how context moved across the org (think standard project channels, shared team drives, and mandated artifacts as part of the product development process). That built a corpus we feed our company brain, and everyone taps it to make better decisions about their own work.

WHY A SPOTIFY TOOL

I wanted the team to feel it, not take my word for it.

The company brain worked because people with more context make better decisions. But even the best process requires change management. So I created a visceral example and let everyone hear the difference on something familiar: their music.

Better memory → better outcomes.

Spotify's API remembers 50 plays.

Spotify remembers everything.

A new app gets ~50 recent plays plus three rank-only top-track lists. No play counts, no timestamps, no history. Spotify also removed the recommendations endpoint for new apps, so similar-tracks is off the table. But every user can export their full streaming history, and mine goes back to 2019.

50

PLAYS THE API SHOWS

A context window. Recency and rank, nothing else.

56,784

STREAMS MY EXPORT REMEMBERS

Every stream since 2019. 30,799 unique tracks, 7.3 years.

1,135x

THE MEMORY MULTIPLIER

Same account, same music, same API scopes. Only memory differs.

Every AI system has **this exact gap.**

This is the same gap every AI product has. Music just makes it obvious.

◆ The API window

What the system can see right now. Recency and rank, no explanation. Ask it why a song is a favorite and the best it can say is "it's in your top tracks."

= A CONTEXT WINDOW

Better memory → better outcomes.

◆ The export

Play counts, completions, skips, deliberate starts, hour-of-day, loyalty across years. Actual behavior, with signals no cleverness can reconstruct from a 50-play window.

= LONG-TERM MEMORY

◆ The labels


Claude listened through my library once and labeled 403 tracks with emotions. Written down, "melancholy" becomes a query instead of a guess made fresh every session.

= SEMANTIC MEMORY



A memory and scoring layer between Claude and Spotify.

● Claude + Spotify's official connector




Spotify picks the songs on their own servers, using the listening data they keep. You get good picks, but you can't see that data, ask why a song was chosen, or build on it.

I tested it: the same throwback ask, run twice, returned two different playlists. Checked against my export, 13 of its 25 picks are songs I've played fewer than 10 times ever.

Developers building their own apps don't get this engine at all. They get the last 50 plays, and that's the path the next two slides measure.

● Claude + local-mood



Claude interprets the ask. The layer remembers every stream, computes the signals, and makes the picks. The wrapped API just executes.

Both apps can answer this. One answers better.

```
compare_memory("throwback", count=15)
```

Throwback: older tracks you still hold onto. Identical algorithm, run twice. Both are real playlists on my Spotify.

1



Can't Help Falling in Love
Elvis Presley

2



Vegeta - Super Saiyan
Bruce Faulconer

3



Sing For The Moment
Eminem

4



Till I Collapse
Eminem, Nate Dogg

5



Somewhere Over The Rainbow/What A Wonderful World
Israel Kamakawiwo'ole

6



The Chain - 2001 Remaster
Fleetwood Mac

7



Jamming
Bob Marley & The Wailers

8



Streets of Philadelphia - Single Edit
Bruce Springsteen

9



Mo Money Mo Problems (feat. Puff Daddy & Mase)
The Notorious B.I.G., Mase, Diddy

● Run 1 • API only

It saw me play Sweet Caroline and Elvis at one party and decided that's who I am. Two to five plays each, no history behind any of it.

1



Mo Money Mo Problems (feat. Puff Daddy & Mase)
The Notorious B.I.G., Mase, Diddy

2



In the End
Linkin Park

3



All Eyez On Me (ft. Big Syke)
2Pac, Big Syke

4



Gangsta's Paradise
Coolio, L.V.

5



Till I Collapse
Eminem, Nate Dogg

6



The Next Episode
Dr. Dre, Snoop Dogg

7



93 'Til Infinity
Souls Of Mischief

8



Numb / Encore
JAY-Z, Linkin Park

9



Still D.R.E.
Dr. Dre, Snoop Dogg

● Run 2 • with memory

What I actually grew up on: Biggie, Pac, Dre, Linkin Park. 13 of 15 picks changed. Numb / Encore has 322 lifetime plays; the API can't see one of them.

This ask, the API can't answer at all.

`compare_memory("comfort", count=15)`

Comfort: long-loved tracks you finish and rarely skip.

```
// everything the official Web API
// gives a new app
top_tracks      rank order only
recently_played last 50 streams
saved_tracks    yes / no

// what "comfort" is made of
completion ratio not in the API
skip rate        not in the API
years of loyalty  not in the API
```

● Run 1 • API only

The inputs don't exist, so local-mood refuses to fake it: 0 of 15 picks reproducible from the window.

- 1  Yours
SG Lewis
- 2  Before Paris
Tom Misch
- 3  It Runs Through Me
Tom Misch, De La Soul
- 4  Sad People
E Kid Cudi
- 5  Know Myself
E Justine Skye, Vory
- 6  On My Mind (Jorja Smith X Preditah)
Jorja Smith, Preditah
- 7  White
Frank Ocean, John Mayer
- 8  Thinkin Bout You
Frank Ocean
- 9  Swim Good
Frank Ocean

● Run 2 • with memory

On my Spotify since July 21. Every pick is backed by completion, skip, and loyalty numbers from 7.3 years of streams.

Music that fits the day I've had. **One sentence to get it.**

The outcome I wanted: the mood I'm in, the ramp I need next, and a few songs I've never heard. Memory does the work.

ME Given the day I've had across Telegram and Slack, figure out the mood I'm in and build a playlist that walks me from there up to leg-day energy. Sneak in a few songs I've never heard.

CLAUDE telegram · get_briefing("today") ✓ · slack · recap ✓ read: long day, two escalations

mood read: drained · target: energetic, motivated

generate_playlist ✓ 20 tracks, ordered calm to energetic from my semantic labels

discover ✓ my database says energetic for me is 2000s hip-hop, so it runs a Spotify catalog search: genre "hip hop", 2000–2009

✓ ~100 candidates back from Spotify, checked one by one against my database: 3 have never been played, those get mixed in

play(device="Zain's iPhone") ✓ playback started

No recommendation engine anywhere. One Spotify search, plus a database that knows every song I've played.

Cronos Labs now runs **on this loop.**

Capture, persist, recall. This is how we work today.

◆ **Decisions, written down**

P/PS docs and PRDs capture the why behind every call, once. Teams stop re-litigating last quarter's decisions and spend that energy shipping. Decision memory is what makes "less deliberating" possible.

DECISION MEMORY

Better memory → better outcomes.

◆ **Retros write learnings**

Our content machine runs weekly engagement retros that write what worked back into per-channel learnings files. Every campaign draft starts from everything the last one taught us instead of from zero.

LEARNING LOOPS

◆ **Context, pre-computed**

A daily job drops the overnight Telegram briefing into the workspace before anyone sits down. People and agents start the day already caught up, instead of spending the first meeting rebuilding context.

STANDING CONTEXT



Clone it. Drop in your export. **Hear the difference.**

MIT licensed. Works with any MCP client. Request your Spotify export today, drop the files in, and run the same comparison on your own history.

github.com/zainbacchus/local-mood-mcp

